


# Neural Nets 2

## Training & Tuning

## Overview:

- **Loss Functions**
- **Backpropagation**
- **Cross Validation**
- **Over- and Underfitting**
- **Tipps & Tricks (e.g. Learning Rate Schedule)**
- **CNN**

# Compute forward propagation (Prediction)



## Compute error/loss $E(y, \hat{y})$



## 1D Output:

$$E = \frac{1}{2m} \sum_{i=1}^m (\hat{y}_i - y_i)^2$$

## nD Output:

$$E = \frac{1}{2m} \sum_{i=1}^m \frac{1}{n} \sum_{j=1}^n (\hat{y}_{ij} - y_{ij})^2$$

$$E = \frac{1}{2mn} \sum_{i=1}^m \sum_{j=1}^n (\hat{y}_{ij} - y_{ij})^2$$

## Compute error/loss $E(y, \hat{y})$



## 1D Output ( $m$ : Batch Size):

$$E = \frac{1}{2m} \sum_{i=1}^m (\hat{y}_i - y_i)^2$$

## $n$ D Output:

$$E = \frac{1}{2m} \sum_{i=1}^m \frac{1}{n} \sum_{j=1}^n (\hat{y}_{ij} - y_{ij})^2$$

$$E = \frac{1}{2mn} \sum_{i=1}^m \sum_{j=1}^n (\hat{y}_{ij} - y_{ij})^2$$

$\hat{y}$	$y$
0	1
1	0
0	0

$$E = \frac{1}{2mn} \sum_{i=1}^m \sum_{j=1}^n (\hat{y}_{ij} - y_{ij})^2$$

$\hat{y}$	$y$
0	1
1	0
0	0

$\hat{y}$	$y$
0	1
1	0
0	0

$$E = \frac{1}{2mn} \sum_{i=1}^m \sum_{j=1}^n (\hat{y}_{ij} - y_{ij})^2$$

$$m = 1, n = 3$$



$\hat{y}$	$y$
0	1
1	0
0	0

$$E = \frac{1}{2mn} \sum_{i=1}^m \sum_{j=1}^n (\hat{y}_{ij} - y_{ij})^2$$

$$m = 1, n = 3$$
$$E = \frac{1}{2 * 1 * 3} \sum_{i=1}^1 \sum_{j=1}^3 (\hat{y}_{ij} - y_{ij})^2$$

$\hat{y}$	$y$
0	1
1	0
0	0

$$E = \frac{1}{2mn} \sum_{i=1}^m \sum_{j=1}^n (\hat{y}_{ij} - y_{ij})^2$$

$$m = 1, n = 3$$

$$E = \frac{1}{2 * 1 * 3} \sum_{i=1}^1 \sum_{j=1}^3 (\hat{y}_{ij} - y_{ij})^2$$

$$E = \frac{1}{6} \sum_{j=1}^3 (\hat{y}_j - y_j)^2$$

$\hat{y}$	$y$
0	1
1	0
0	0

$$E = \frac{1}{2mn} \sum_{i=1}^m \sum_{j=1}^n (\hat{y}_{ij} - y_{ij})^2$$

$$m = 1, n = 3$$

$$E = \frac{1}{2 * 1 * 3} \sum_{i=1}^1 \sum_{j=1}^3 (\hat{y}_{ij} - y_{ij})^2$$

$$E = \frac{1}{6} \sum_{j=1}^3 (\hat{y}_j - y_j)^2$$

$$E = \frac{1}{6} ((0 - 1)^2 + (1 - 0)^2 + (0 - 0)^2) = \frac{1}{3}$$

## Logarithmic Loss (Cross Entropy)

- Captures the intuition of classification
- Can be used for output probabilities (**e.g. Softmax**)
- $p$ : Predicted probability of a class (confidence)
- **Binary:**  $E = -(y * \log(p) + (1 - y) * \log(1 - p))$
- **Multiclass:**  $E = -\sum_{j=1}^n y_j * \log(p_j)$

$$E = -(y * \log(p) + (1 - y) * \log(1 - p))$$

**Assume**  $y = 1$ :


$$E(y = 1, p) = -(1 * \log(p) + (1 - 1) * \log(1 - p))$$

$$E(y = 1, p) = -(1 * \log(p))$$

$$E(y = 1, p) = -\log(p)$$

$$E = -(y * \log(p) + (1 - y) * \log(1 - p))$$

**Assume  $y = 1$ :  $E(y = 1, p) = -\log(p)$**




$$E = - \sum_{j=1}^n y_j * \log(p_j)$$

$$E = -(1 * \log(0.1) + 0 * \log(0.8) + 0 * \log(0.1))$$

$$E = -\log(0.1)$$

$\hat{y}$	$y$
0.1	1
0.8	0
0.1	0




$$E = - \sum_{j=1}^n y_j * \log(p_j)$$

$$E = -(1 * \log(0.8) + 0 * \log(0.1) + 0 * \log(0.1))$$

$$E = -\log(0.8)$$



$\hat{y}$	$y$
0.8	1
0.1	0
0.1	0





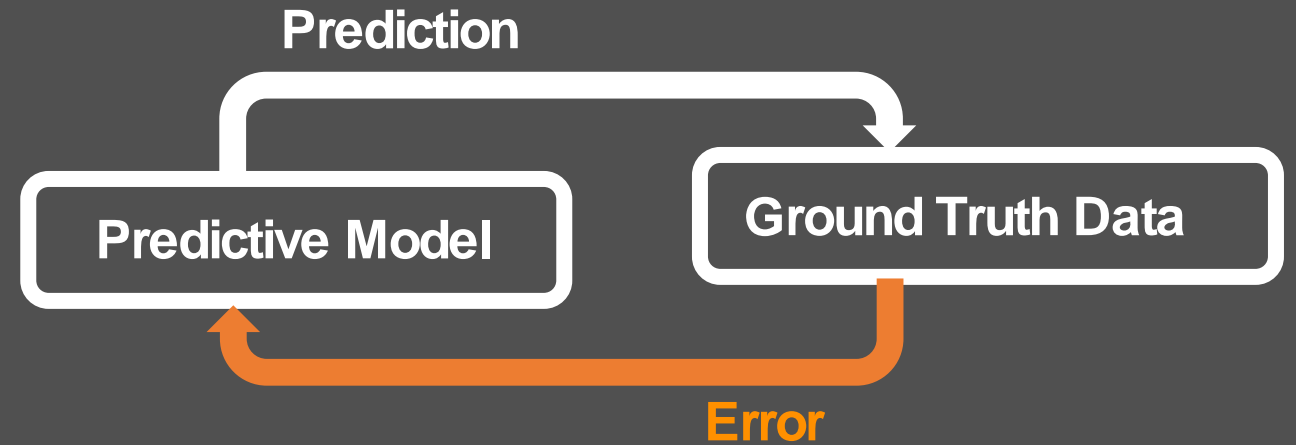
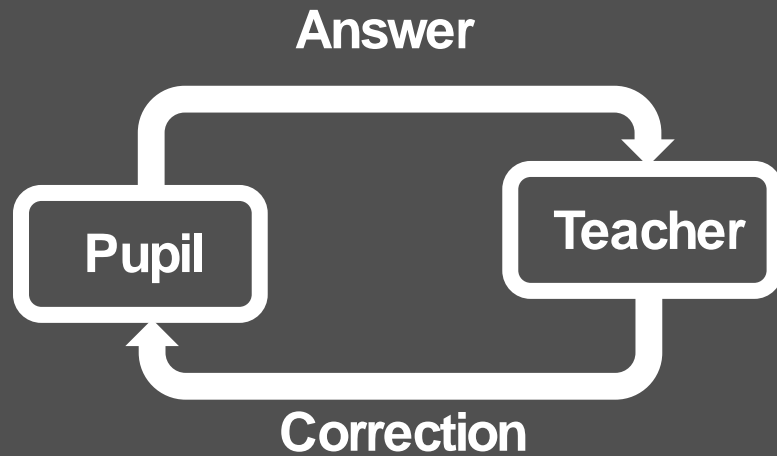
## Two Processes:

- Forward Propagation (Prediction)
- **Backward Propagation (Weight Tuning)**



# Backpropagation

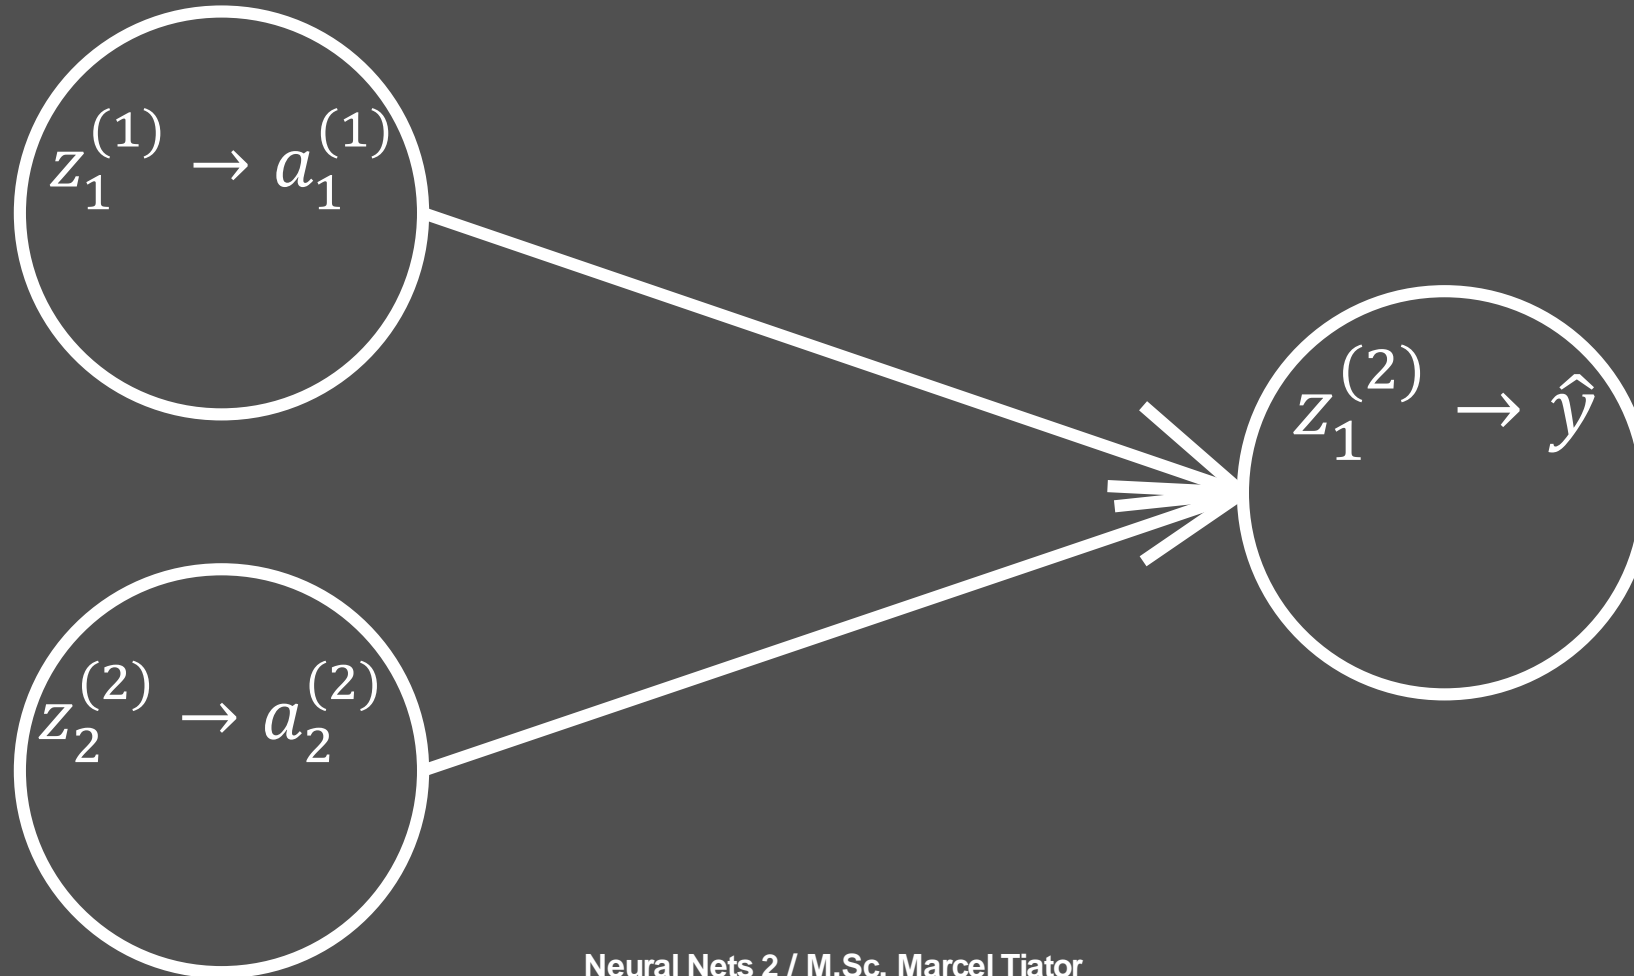
- Compare to more complex Gradient Descent
- Training of Neural Net



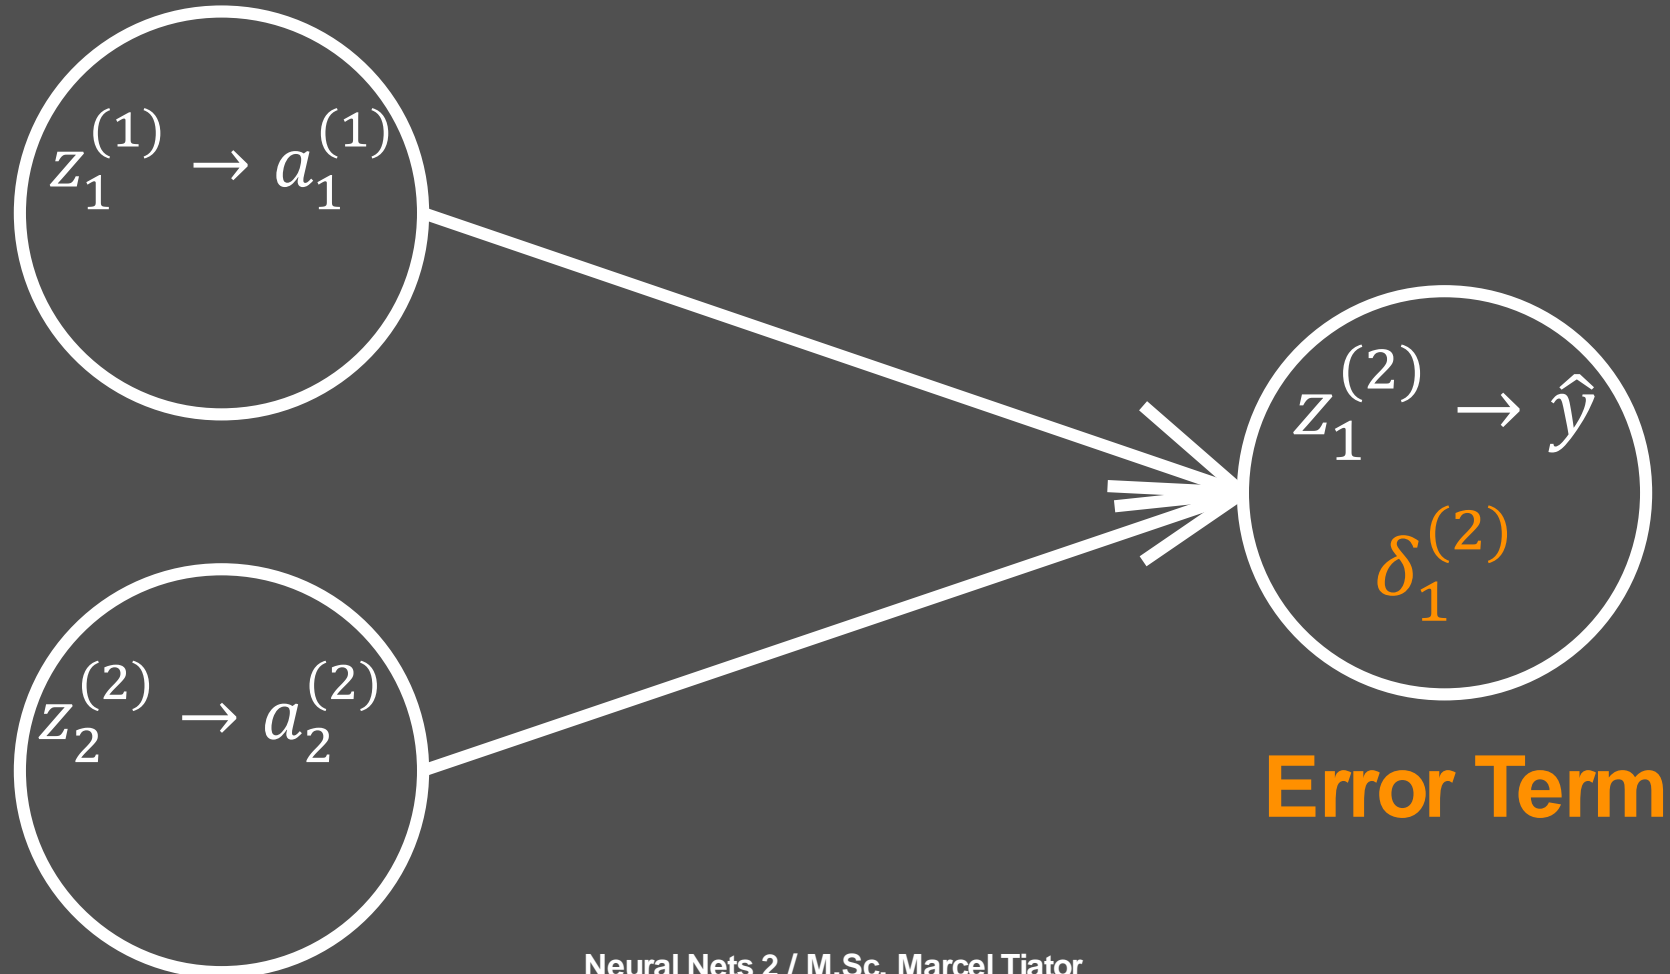
## Algorithm (Given input $x$ )

- **Compute forward propagation**
  - Get a prediction  $\hat{y}$
- **Compute error/loss  $E(y, \hat{y})$**
- **Propagate the error back through the network**
  - Tune weights

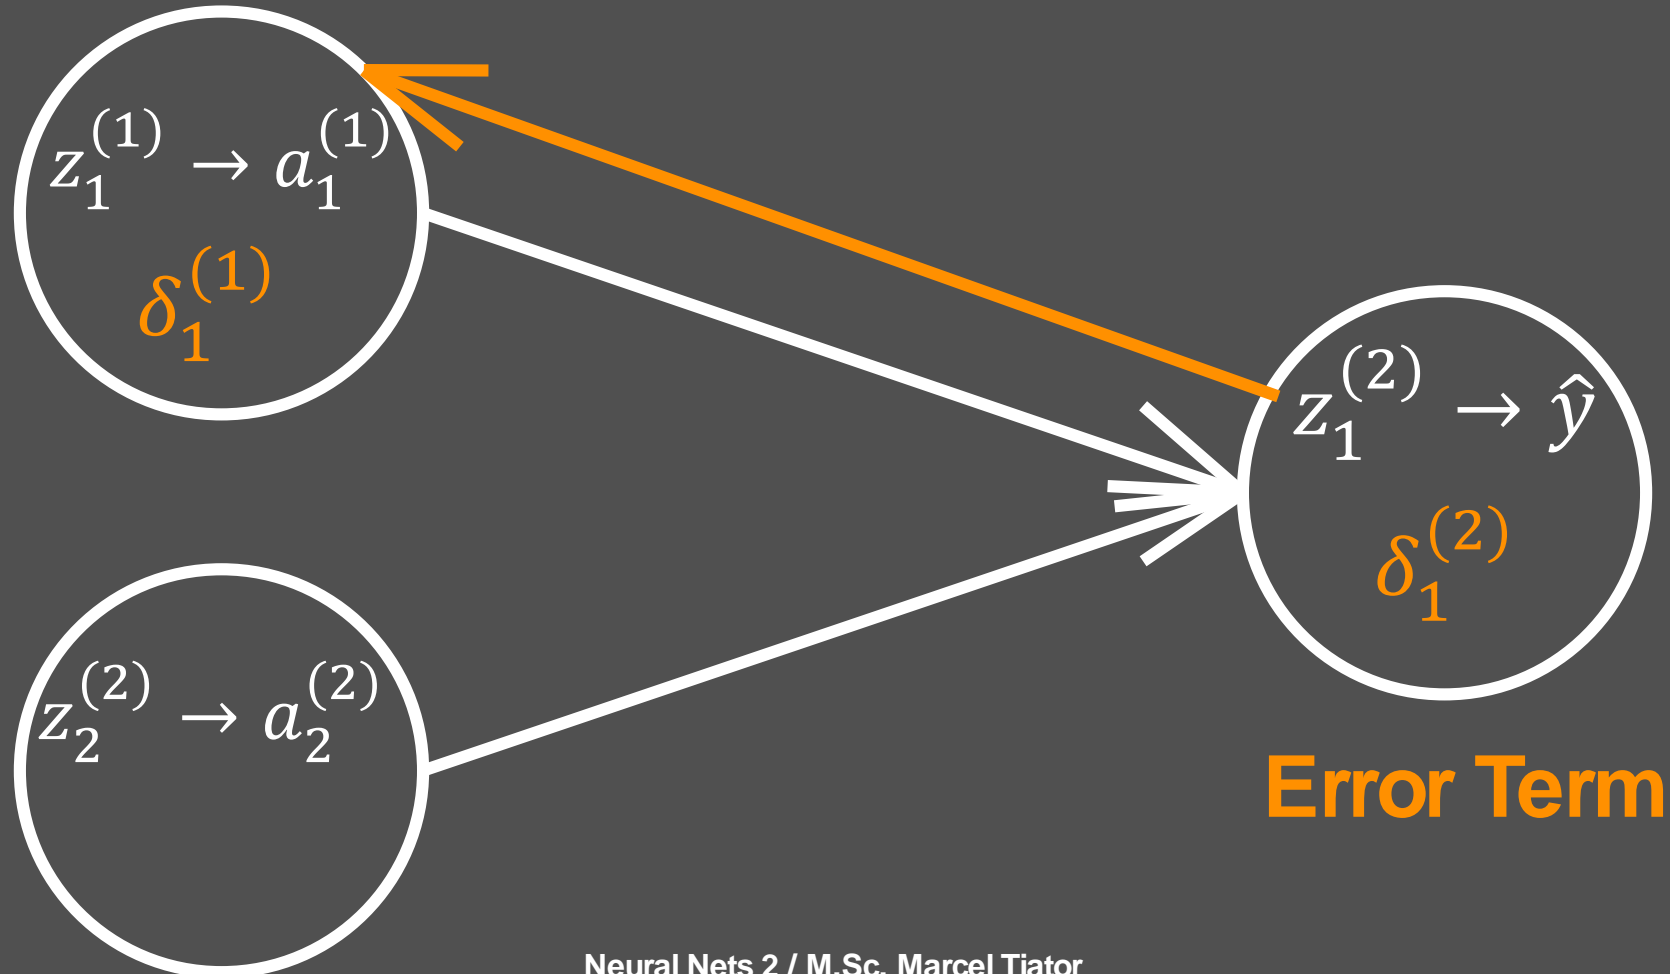
## Propagate the error back through the network




## Propagate the error back through the network



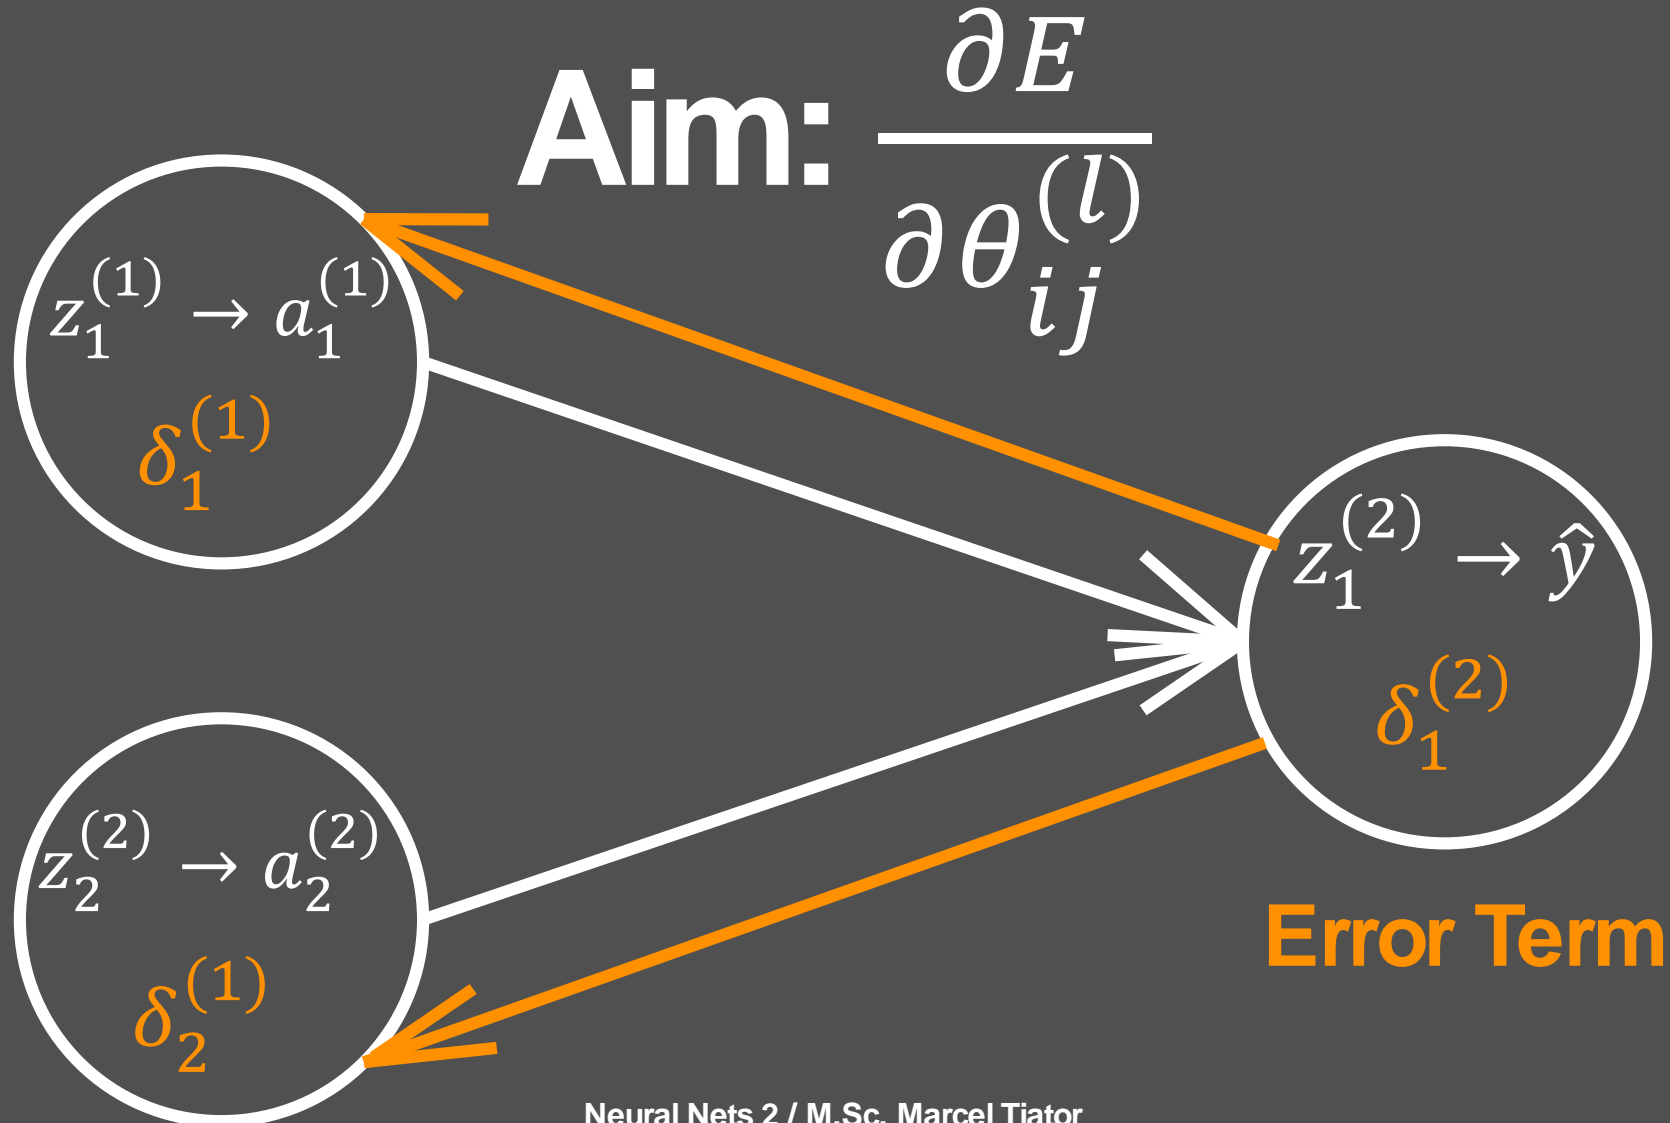
## Propagate the error back through the network



## Propagate the error back through the network



# Backpropagation






## Parameterupdate:

$$\theta_{ij}^{(l)} := \theta_{ij}^{(l)} - \alpha \frac{\partial E}{\partial \theta_{ij}^{(l)}}$$

(Remember **Gradient Descent**)



# Number of weights in MLP

$L$ : Number of layer

$s_l$ : Number of neuron in layer  $l$

$$\#weights = \sum_{l=1}^{L-1} s_{l+1} * (s_l + 1)$$

**Bias-Unit ( $x_0, a_0^{(l)}$ )**



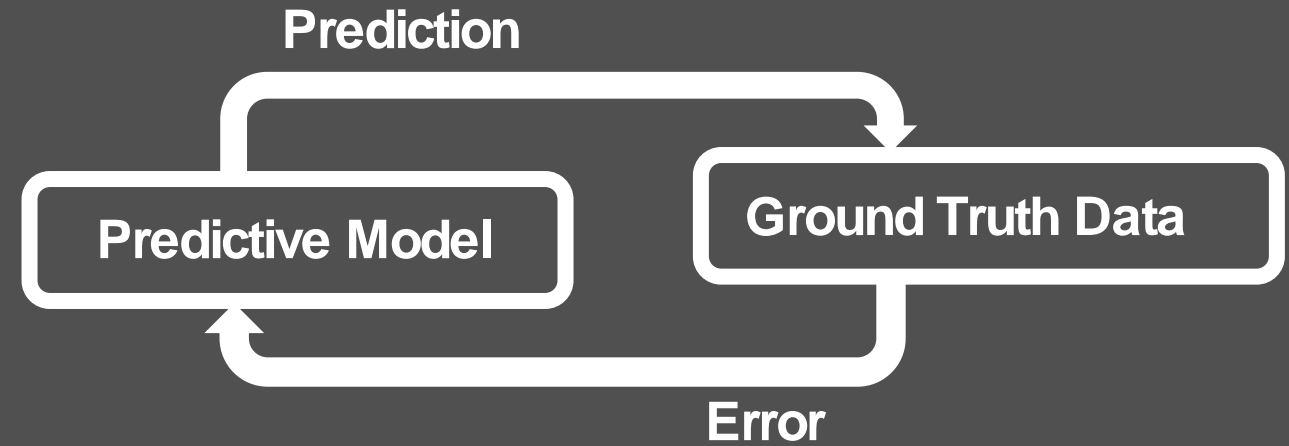
**Given a MLP with following architecture: [3,4,4,3]**  
**Calculate the *#weights* of the MLP?**

$$\#weights = \sum_{l=1}^{L-1} s_{l+1} * (s_l + 1)$$

$$\#weights = (4 * (3 + 1)) + (4 * (4 + 1)) + (3 * (4 + 1))$$
$$\#weights = 16 + 20 + 15 = 51$$


# Learning Curves


- **x-Axis: Training Step**
- **y-Axis: Error/Loss**




# Rate performance of neural net


- High bias / underfitting
- High variance / overfitting




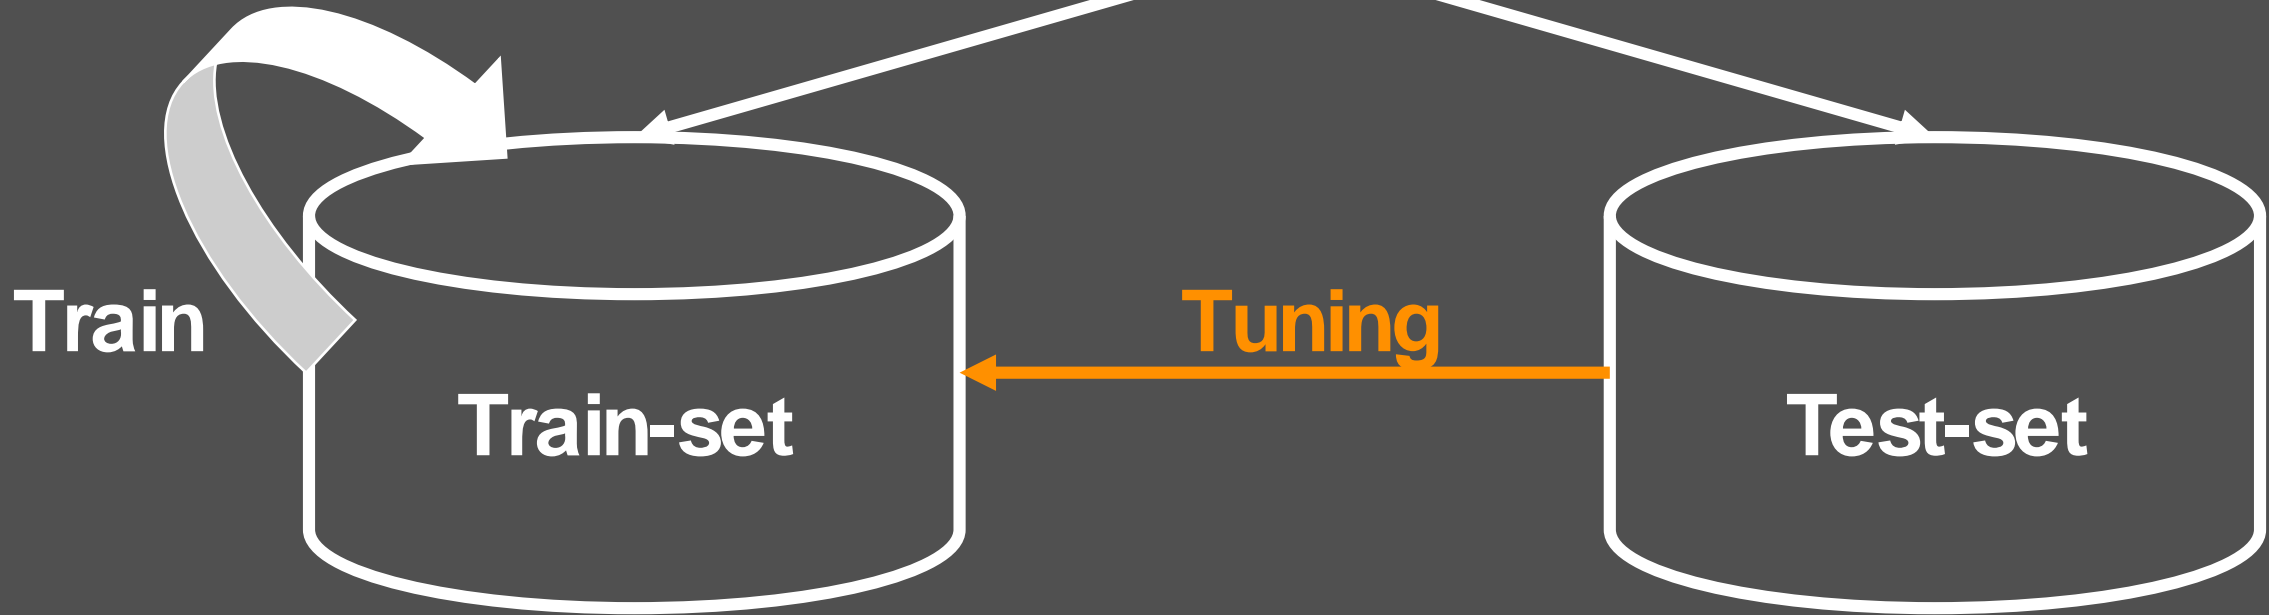














# Cross Validation

- Tune your network with test-set performance (e.g. accuracy)
- Attention: Overfitting of test-set is possible
- Try 80% of data as train-set



# Cross Validation

- Real Pro's use an additional validation-set
- Prevent overfitting
- Validation-set is a third separate dataset
- Try 60% Train, 20% Test, 20% Validation



## High Bias / Underfitting

- **High Training Error & High Test Error**
- **Complexity of neural net is too low (Too less weights)**
- **Add layer or units (neurons in layer)**
- **If possible: add more features of data**

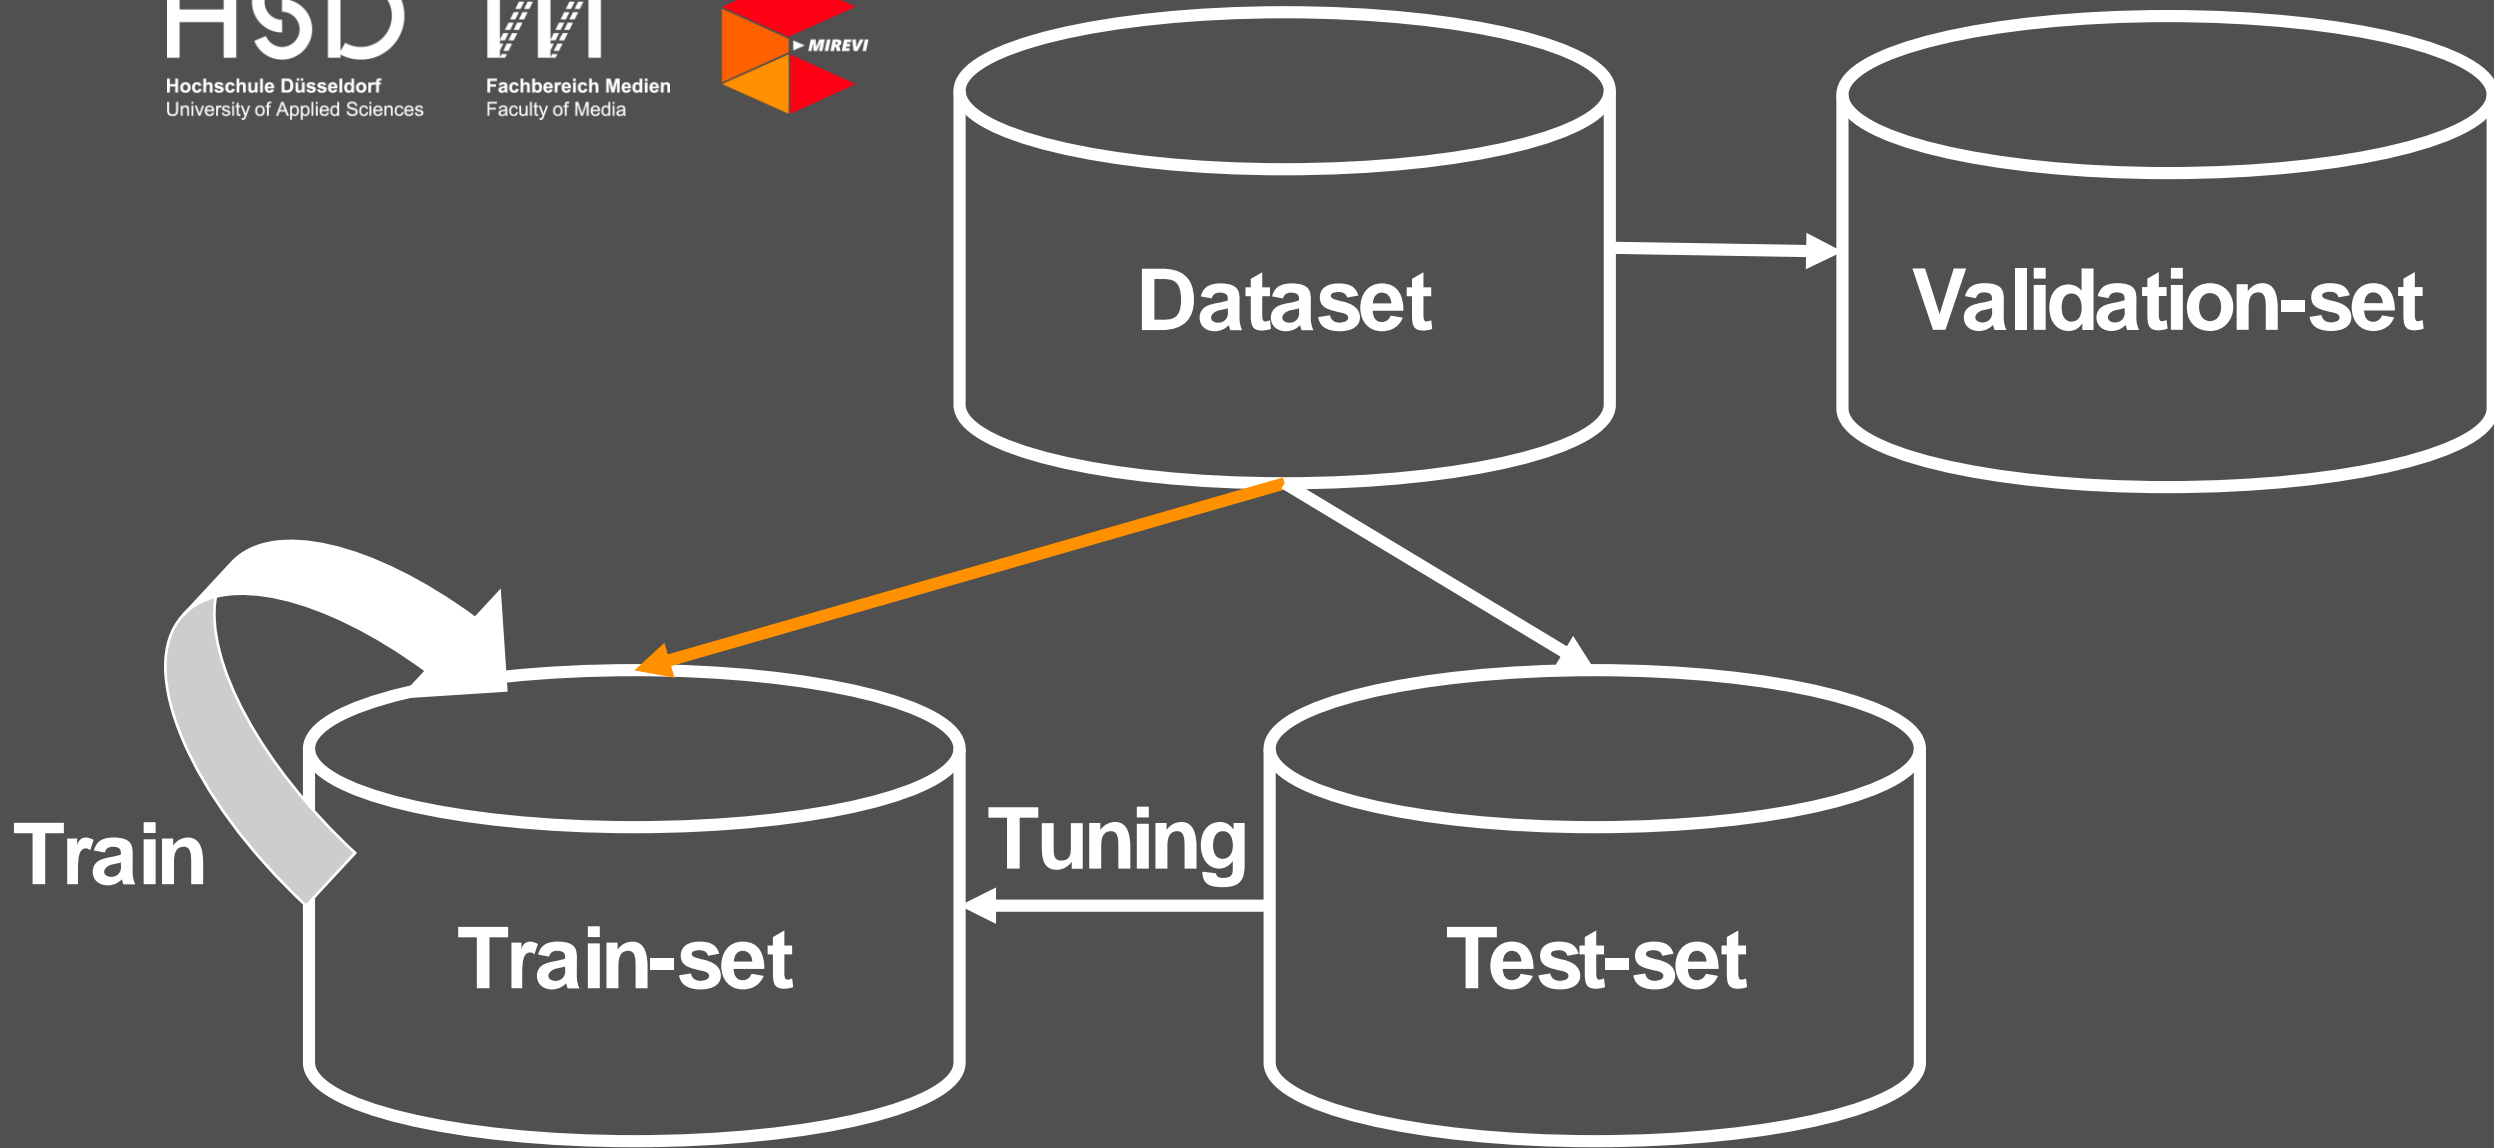
## High Bias / Underfitting

- **Create polynomial features**
  - Add column with  $x_1^2, x_2^2, x_1x_2, \dots$
- **More trainings-examples will have no significant effect**
- **Reduce regularization**

## High Variance / Overfitting

- Neural net has learned the noise of the data
- Lower complexity
- More training examples can help
- Use regularization





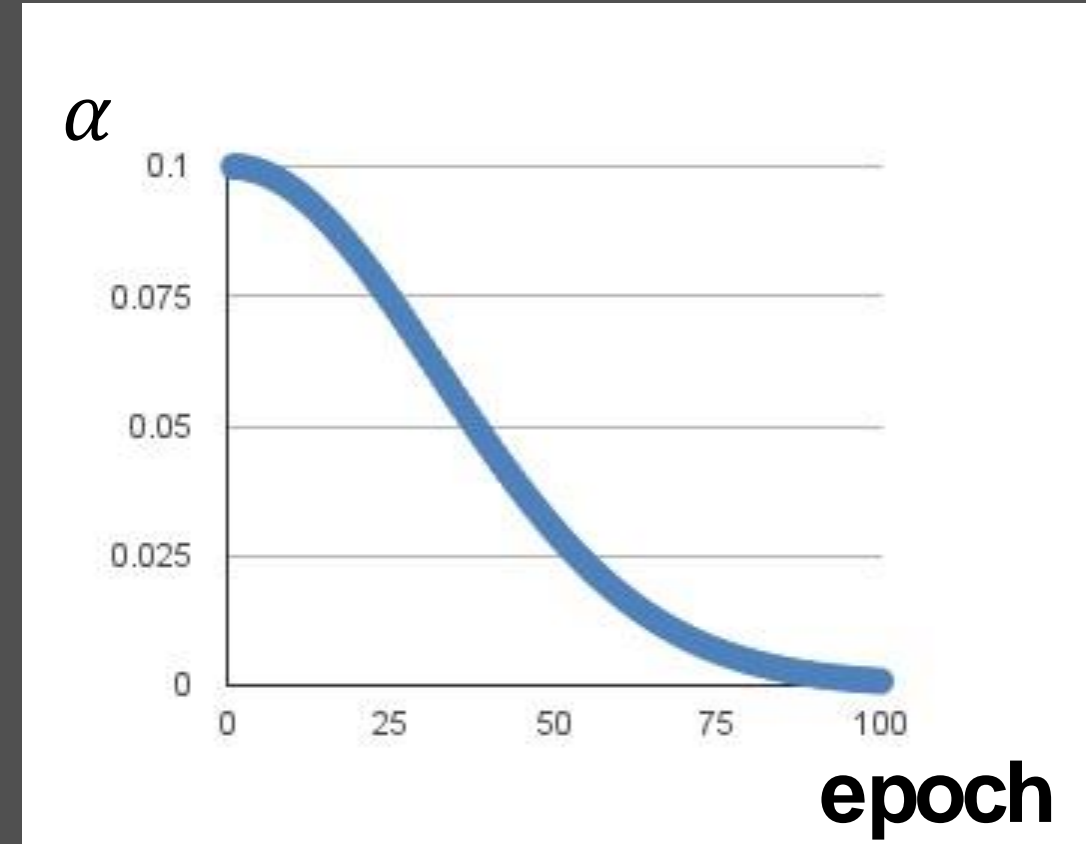
## Regularization:

- Try to keep the values of the weights small
- $\lambda$ : Regularization parameter
- **L1-Regularization:**  $E = f(\hat{y}, y) + \frac{\lambda}{m} \sum |\theta|$
- **L2-Regularization:**  $E = f(\hat{y}, y) + \frac{\lambda}{m} \sum \theta^2$


# Please scale your data!

# Learning Rate Decay (Learning Rate Schedule)

$$\alpha_n = \frac{\alpha}{1 + decay * epoch}$$

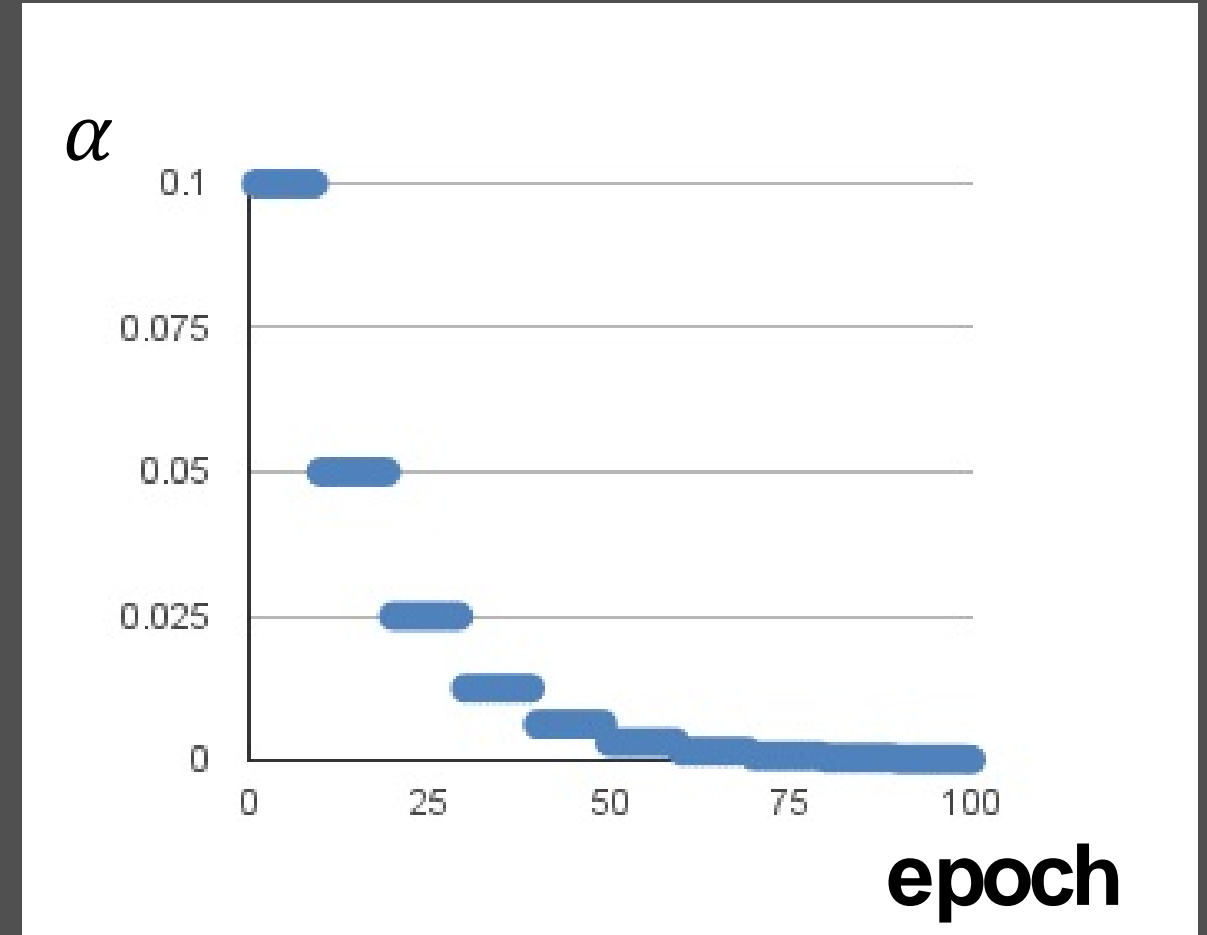


Prevent this  
with learning  
rate schedule



# Drop Based Learning Rate Schedule


e.g. Half the learning rate every 10th epoch



# Convolutional Neural Nets (CNN)

# Disadvantage of using MLP in image recognition

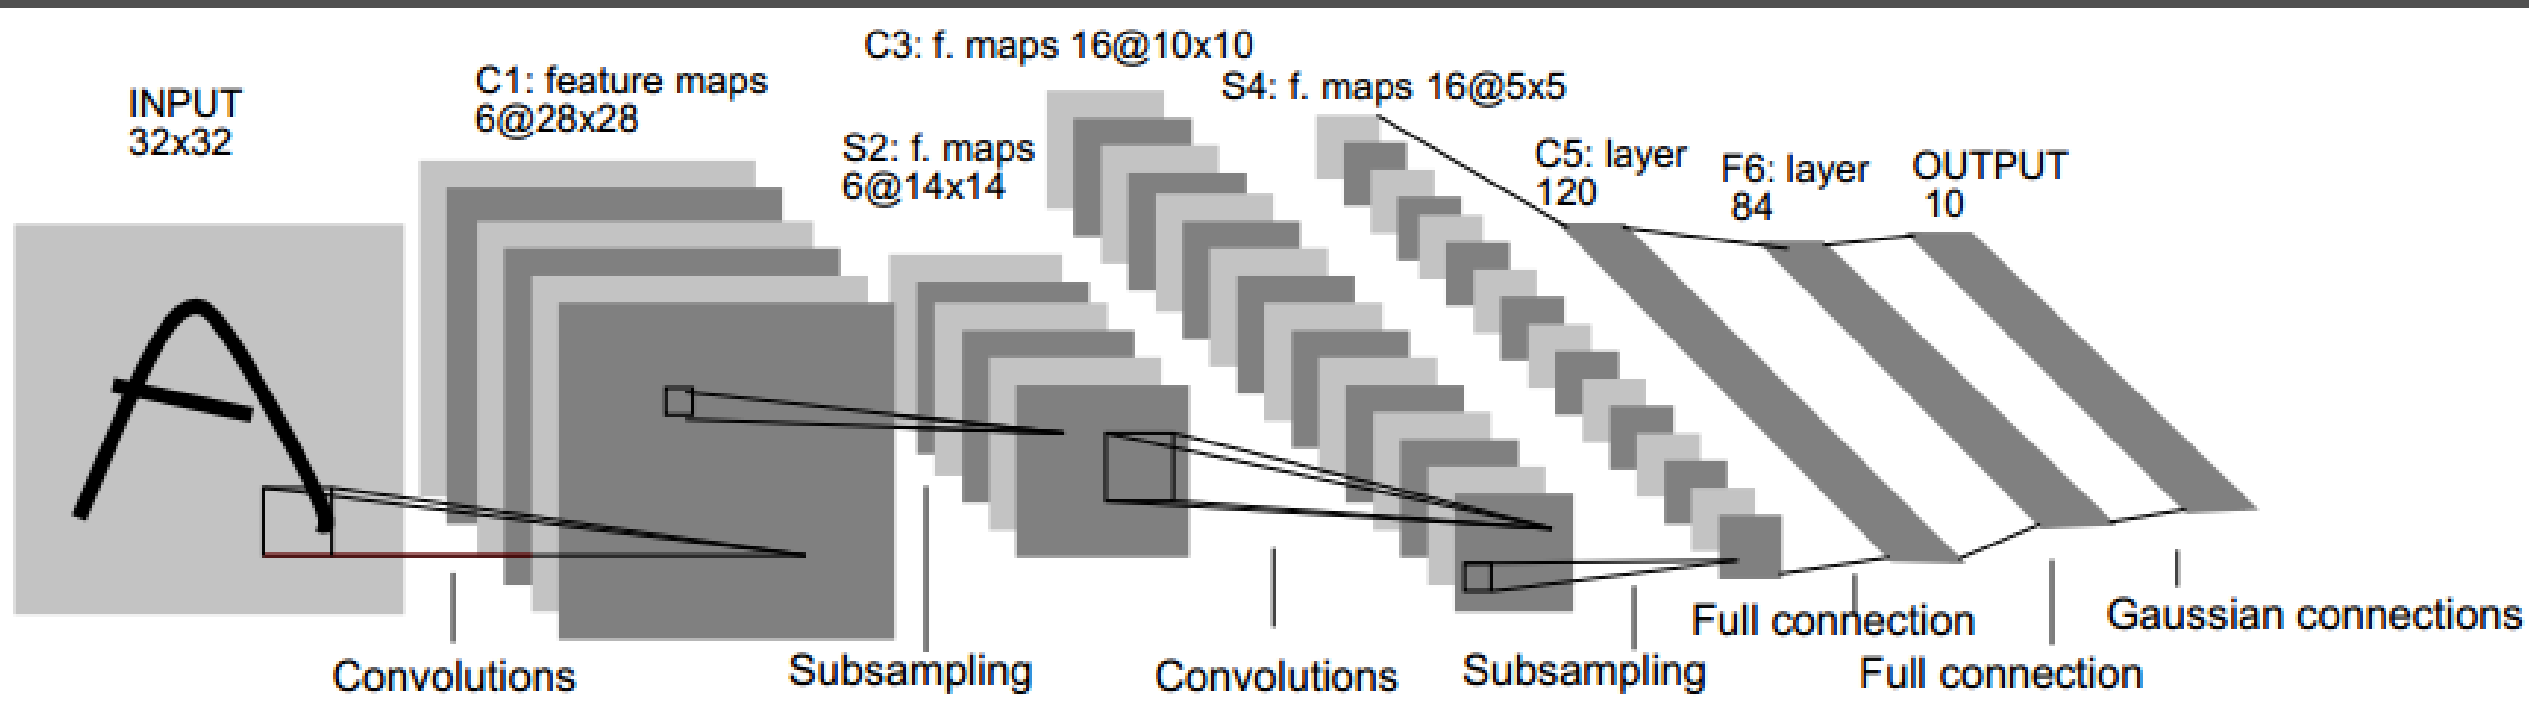
- You can fool them by translating the object





# Convolutional Neural Nets (CNN's)


- 2D CNN's assume **images** as input
- They can learn **transformational invariance**
- Different architecture than MLP
- Trained with Backpropagation



# Architectural elements of a CNN

- Convolutional Layer
- Max-Pooling Layer
- Fully Connected Layer

# CONVOLUTIONS




# Convolutional Layer

- Apply image filters
- Feature detectors in combination with activation
- Special parameters:
  - Stride
  - Padding

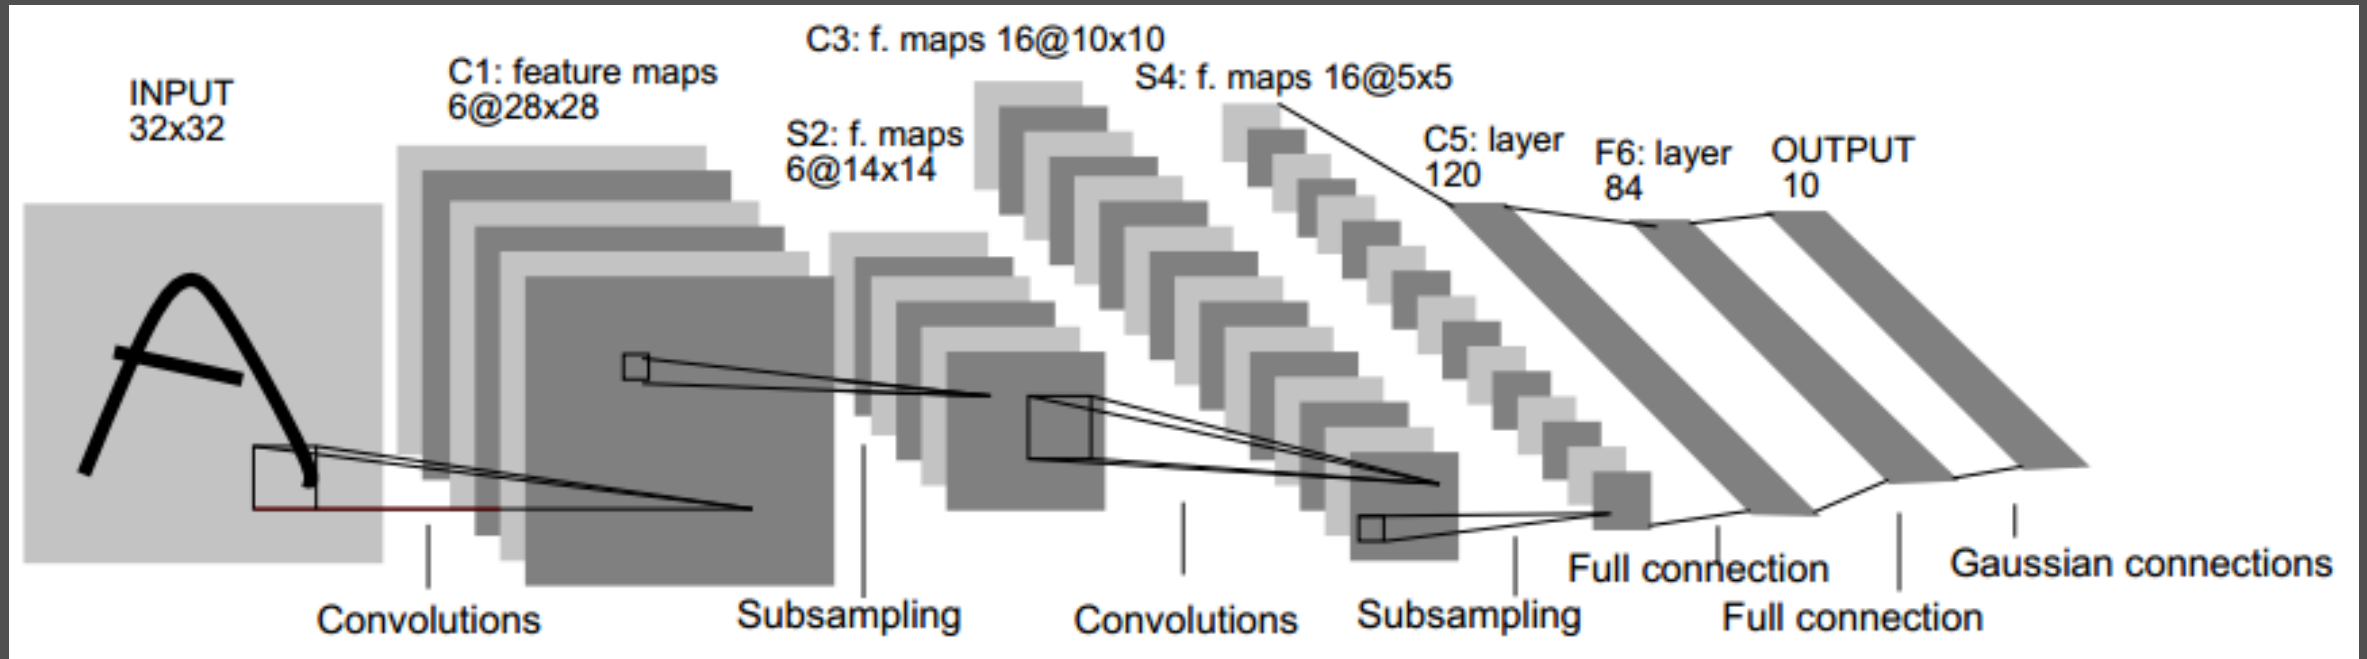
## Activation after convolution

- We will apply an activation function after convolution
- Modern CNN's use ReLu




# Max-Pooling Layer

- Dimensionality reduction
- Pass the greatest value of a field to the next layer




# Fully Connected Layer

- Flatten the 2D (in case of image) to 1D layer
- Use fully connected layer up to the output





# Sometimes useful: Normalization Layer



# Star Recognition



## Today's exercise:

- Complete the code to train a CNN
- Real world image classification problem
- Given a face of a person: classify star

## Sources:

- `cnn.py`
- `cnn_train.py`
- `cnn_eval.py`
- `cnn_predict_cv.py`
- `download_image_data.py`
- `write_tfrecord.py`
- `read_tfrecord.py`
- `utils.py`

## How to get image data?

- Execute the steps in the README.md of this project
- Execute download\_data.py
- Faces will be extracted through haarcascade detection (OpenCV)
- **Already done**, if you download the data from the Nextcloud

## utils.py (**First Step**):

- Mainly image processing
- Face Extraction
- Grayscale
- Scaling to quadratic size
- Normalization

Write .tfrecord file (will be extended soon, **Second step**)

- **Execute** the **write\_tfrecord.py** script
- Special binary-dataset will be written to disk
- We will directly apply **image preprocessing**
  - Execute read\_tfrecord.py to validate the contents of the .tfrecord file (**Third Step**)
- Execute **read\_tfrecord.py** to **validate** the **contents** of the .tfrecord file

# cnn.py – Convolutional Neural Network

- Definition of the architecture
- Loss function
- Definition of training and optimizer
- **Script not executable**



## `cnn_train.py` – training of the CNN (**Fourth Step**)

- Executable script
- Tuning of hyperparameters (e.g. learning rate, etc.) possible
- Read .tfrecord file
- Manages the training process

## `cnn_eval.py` (**Sixth Step**)

- Executable
- Possibility to validate the model with a test-set

## `cnn_predict.py` (**Seventh Step**)

- Executable
- Prediction of star with one image

## Tasks:

- Image Processing in `utils.py`
- Write `.tfrecord` file with `write_tfrecord.py`
- Validate `.tfrecord` with `read_tfrecord.py`
  - Tipp: If you leave out preprocessing (without float conversion and normalization) of image you should see a colored image of the face of a star
- Prepare training of CNN in `cnn_train.py`
- Evaluate the result with `cnn_eval.py`
- Classify one image with `cnn_predict_cv.py`
- Add train, test, validation split functionality

- **Data (if not yet downloaded):**  
<https://nextcloud.mirevi.medien.hs-duesseldorf.de/index.php/s/kPXwJiac7vTQVeu>
- **Pretrained Weights:**  
<https://nextcloud.mirevi.medien.hs-duesseldorf.de/index.php/s/L6Y6tnD3PpANKmr>
- **Repository:** <https://github.com/mati3230/modalg181>
- **Read:** [https://www.tensorflow.org/tutorials/deep\\_cnn](https://www.tensorflow.org/tutorials/deep_cnn)