


Neural Nets 3

Performance & Optimization

Overview:

- **CNN Visualization**
- **Performance Metrics**
- **Batch, Mini-Batch and Stochastic Learning**
- **K-Fold-Cross Validation**
- **Momentum**
- **Dropout**
- **Plot while Training in tensorflow**



Visualizing the process within a CNN:


<https://cs.stanford.edu/people/karpathy/convnetjs/demo/cifar10.html>

relu (32x32x16)
max activation: 1.1251, min: 0
max gradient: 0.05035, min: -0.04651

Activations:



Activation Gradients:

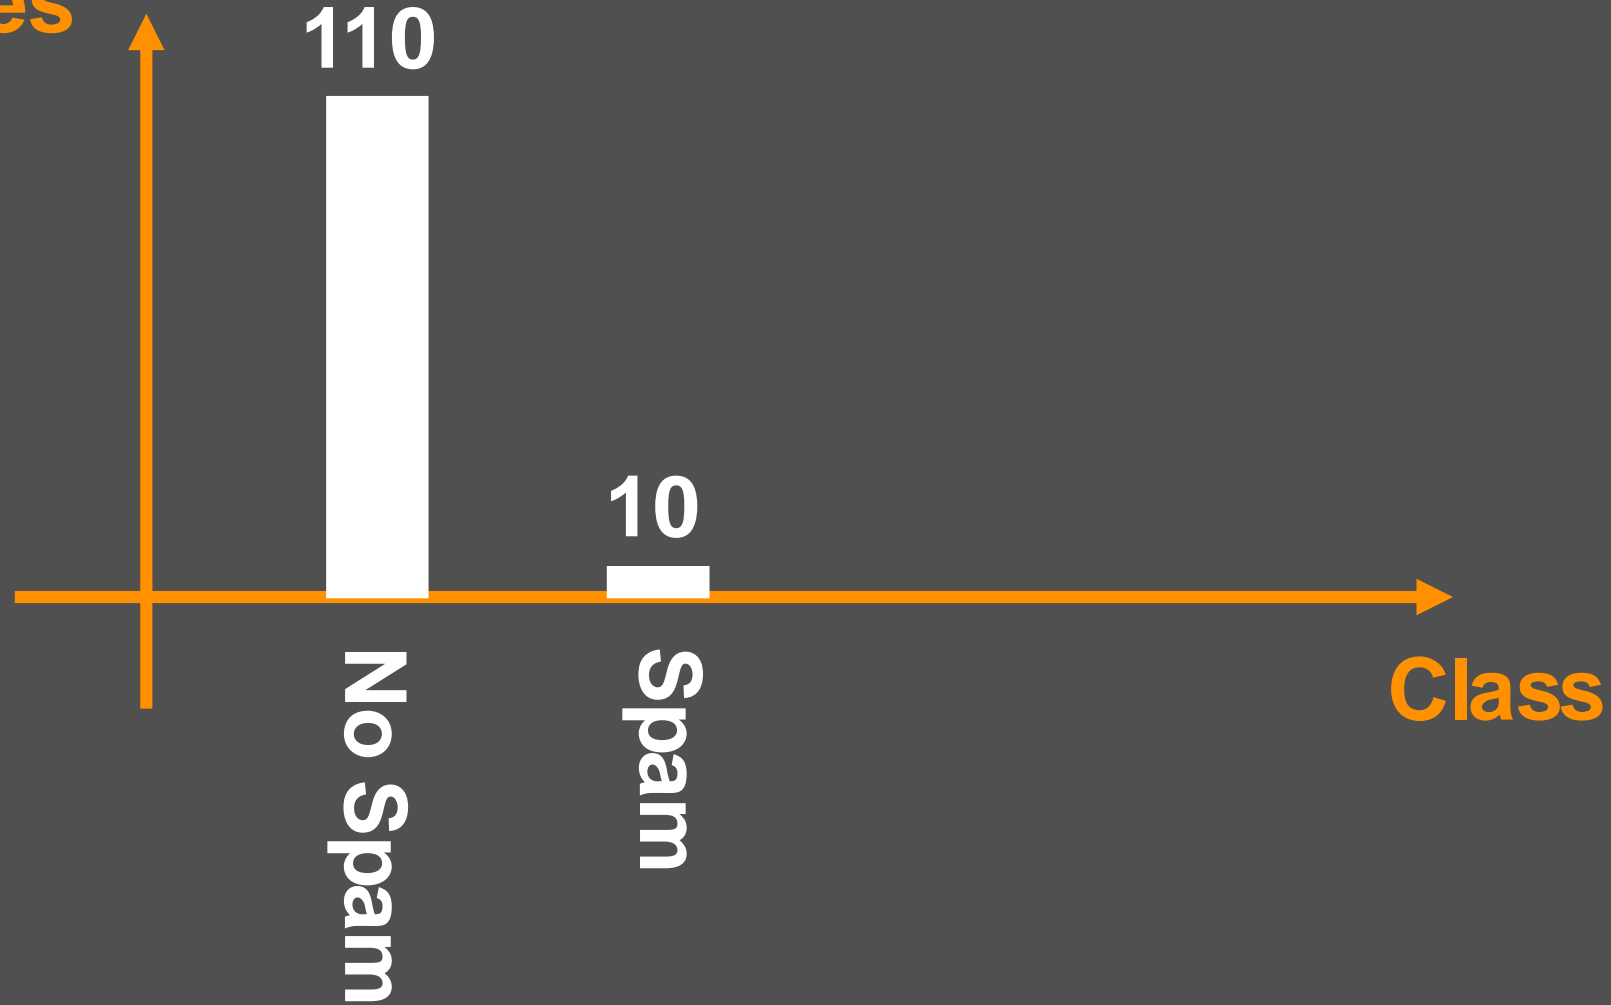


Improve generalization of CNN through **augmentation of dataset:**


- Rotate
- Add Noise
- Mirror
- ...

In which situations is the accuracy of a classifier not an appropriate metric?

#Examples



#Examples



A classifier which always says „NoSpam“ will have an accuracy of $\frac{110}{120} * 100\% = 91.67\%$

More Metrics

True Class	Predicted Class	Type
0	0	True Positive
0	1	False Negative
1	0	False Positive
1	1	True Negative

Confusion Matrix

$$\begin{matrix}
 & \begin{matrix} 0 & 1 \end{matrix} \\
 \begin{matrix} 0 \\ 1 \end{matrix} & \begin{pmatrix} \#True\ Positive & \#False\ Negative \\ \#False\ Positive & \#True\ Negative \end{pmatrix}
 \end{matrix}$$


True Class	Predicted Class	Type
0	0	True Positive
0	1	False Negative
1	0	False Positive
1	1	True Negative

$$\begin{pmatrix} \#True\ Positive & \#False\ Negative \\ \#False\ Positive & \#True\ Negative \end{pmatrix}$$

$$Accuracy = \frac{\#True\ Positive + \#True\ Negative}{\#Examples}$$

$$\begin{pmatrix} \#True\ Positive & \#False\ Negative \\ \#False\ Positive & \#True\ Negative \end{pmatrix}$$

#Examples



True
 No Spam
 Spam

Predicted
 No Spam
 Spam

$$\begin{pmatrix} 110 & 0 \\ 10 & 0 \end{pmatrix}$$

$$\begin{pmatrix} \#True\ Positive & \#False\ Negative \\ \#False\ Positive & \#True\ Negative \end{pmatrix}$$

	Predicted	
	No Spam	Spam
True	No Spam	110 0
	Spam	10 0

$$Accuracy = \frac{\#True\ Positive + \#True\ Negative}{\#Examples}$$

$$Accuracy * 100\% = \frac{110 + 0}{120} * 100\% = 91.67\%$$

$$\begin{pmatrix} \#True\ Positive & \#False\ Negative \\ \#False\ Positive & \#True\ Negative \end{pmatrix}$$

$$Precision = \frac{\#True\ Positive}{\#True\ Positive + \#False\ Positive}$$

$$\begin{pmatrix} \#True\ Positive & \#False\ Negative \\ \#False\ Positive & \#True\ Negative \end{pmatrix}$$

	Predicted	
	No Spam	Spam
True	No Spam	110 0
	Spam	10 0

$$Precision = \frac{\#True\ Positive}{\#True\ Positive + False\ Positive}$$

$$Precision * 100\% = \frac{110}{110 + 10} * 100\% = 91.67\%$$

$$\begin{pmatrix} \#True\ Positive & \#False\ Negative \\ \#False\ Positive & \#True\ Negative \end{pmatrix}$$

$$Recall = \frac{\#True\ Positive}{\#True\ Positive + \#False\ Negative}$$

$$\begin{pmatrix} \#True\ Positive & \#False\ Negative \\ \#False\ Positive & \#True\ Negative \end{pmatrix}$$

	Predicted	
	No Spam	Spam
True	No Spam	110
	Spam	10

$$Recall = \frac{\#True\ Positive}{\#True\ Positive + \#False\ Negative}$$

$$Precision * 100\% = \frac{110}{110 + 0} * 100\% = 100\%$$

$$\textit{Precision} = \frac{\#True\ Positive}{\#True\ Positive + \#False\ Positive}$$

$$\textit{Recall} = \frac{\#True\ Positive}{\#True\ Positive + \#False\ Negative}$$

$$\textit{F1Score} = \frac{2 * \textit{Precision} * \textit{Recall}}{\textit{Precision} + \textit{Recall}}$$

Precision = 1

Recall = 1

$$F1Score = \frac{2 * 0.9167 * 1}{1 + 0.9167} \sim 0.957$$

		Predicted	
		No Spam	Spam
True	No Spam	110	0
	Spam	10	0

$$F1Score = \frac{2 * Precision * Recall}{Precision + Recall}$$

		predicted condition			
total population		prediction positive	prediction negative	Prevalence $= \frac{\Sigma \text{condition positive}}{\Sigma \text{total population}}$	
true condition	condition positive	True Positive (TP)	False Negative (FN) (type II error)	True Positive Rate (TPR), Sensitivity, Recall, Probability of Detection $= \frac{\Sigma \text{TP}}{\Sigma \text{condition positive}}$	False Negative Rate (FNR), Miss Rate $= \frac{\Sigma \text{FN}}{\Sigma \text{condition positive}}$
	condition negative	False Positive (FP) (Type I error)	True Negative (TN)	False Positive Rate (FPR), Fall-out, Probability of False Alarm $= \frac{\Sigma \text{FP}}{\Sigma \text{condition negative}}$	True Negative Rate (TNR), Specificity (SPC) $= \frac{\Sigma \text{TN}}{\Sigma \text{condition negative}}$
		Positive Predictive Value (PPV), Precision $= \frac{\Sigma \text{TP}}{\Sigma \text{prediction positive}}$	False Omission Rate (FOR) $= \frac{\Sigma \text{FN}}{\Sigma \text{prediction negative}}$	Positive Likelihood Ratio (LR+) $= \frac{\text{TPR}}{\text{FPR}}$	Diagnostic Odds Ratio (DOR) $= \frac{\text{LR+}}{\text{LR-}}$
		False Discovery Rate (FDR) $= \frac{\Sigma \text{FP}}{\Sigma \text{prediction positive}}$	Negative Predictive Value (NPV) $= \frac{\Sigma \text{TN}}{\Sigma \text{prediction negative}}$	Negative Likelihood Ratio (LR-) $= \frac{\text{FNR}}{\text{TNR}}$	

$\begin{pmatrix} \#True\ Positive & \#False\ Negative \\ \#False\ Positive & \#True\ Negative \end{pmatrix}$

	Predicted		
	No Spam	Spam	
True	No Spam	3	3
	Spam	5	16

#Examples

= #True Positive + #False Negative + #False Positive + #True Negative

$$\#Examples = 3 + 3 + 5 + 16 = 27$$

$\#Examples = 27$

$$Accuracy = \frac{\#True\ Positive + \#True\ Negative}{\#Examples}$$

$$Accuracy = \frac{3 + 16}{27} \sim 0.7$$

	Predicted	
True	No Spam	Spam
	$\begin{pmatrix} 3 & 3 \\ 5 & 16 \end{pmatrix}$	

$\begin{pmatrix} \#True\ Positive & \#False\ Negative \\ \#False\ Positive & \#True\ Negative \end{pmatrix}$

$$\text{Precision} = \frac{\#True\ Positive}{\#True\ Positive + \#False\ Positive}$$

$$\text{Precision} = \frac{3}{3 + 5} \sim 0.375$$

	Predicted	
	No Spam	Spam
True No Spam	3	3
Spam	5	16

$$\begin{pmatrix} \#True\ Positive & \#False\ Negative \\ \#False\ Positive & \#True\ Negative \end{pmatrix}$$

$$Recall = \frac{\#True\ Positive}{\#True\ Positive + \#False\ Negative}$$

$$Recall = \frac{3}{3 + 3} = 0.5$$

	Predicted		
	No Spam	Spam	
True	No Spam	3	3
	Spam	5	16

$$\begin{pmatrix} \#True\ Positive & \#False\ Negative \\ \#False\ Positive & \#True\ Negative \end{pmatrix}$$

$$F1Score = \frac{2 * Precision * Recall}{Precision + Recall}$$

$$Precision = 0.375$$

$$Recall = 0.5$$

$$F1Score = \frac{2 * 0.375 * 0.5}{0.375 + 0.5} \sim 0.429$$

	Predicted	
True	No Spam	Spam
	<div>3</div> <div>5</div>	<div>3</div> <div>16</div>

$$\begin{pmatrix} \#True\ Positive & \#False\ Negative \\ \#False\ Positive & \#True\ Negative \end{pmatrix}$$

Accuracy = 0.7
Precision = 0.375
Recall = 0.5
F1Score = 0.429

	Predicted	
True	No Spam	Spam
	$\begin{pmatrix} 3 & 3 \\ 5 & 16 \end{pmatrix}$	

$\begin{pmatrix} \#True\ Positive & \#False\ Negative \\ \#False\ Positive & \#True\ Negative \end{pmatrix}$

Given the following Confusion Matrix – calculate:

- Accuracy
- Average Precision
- Average Recall
- F1Score of averaged Precision and Recall

True \ Predicted			
	Dog	Cat	Mouse
	1	0	1
	0	3	2
	Dog	Cat	Mouse
	2	1	2

$$\text{Accuracy} = \frac{\#True\ Positive + \#True\ Negative}{\#Examples}$$

$$\text{Precision} = \frac{\#True\ Positive}{\#True\ Positive + \#False\ Positive}$$

$$\text{Recall} = \frac{\#True\ Positive}{\#True\ Positive + \#False\ Negative}$$

$$F1Score = \frac{2 * Precision * Recall}{Precision + Recall}$$

True \ Predicted	Predicted		
	Dog	Cat	Mouse
Dog	1	0	1
Cat	0	3	2
Mouse	2	1	2

- $\#Example = 12$

True \ Predicted	Predicted		
	Dog	Cat	Mouse
Dog	1	0	1
Cat	0	3	2
Mouse	2	1	2

- $\#Example = 12$
- $Accuracy = 0.5$

True	Predicted		
	Dog	Cat	Mouse
Dog	1	0	1
Cat	0	3	2
Mouse	2	1	2

- $\#Example = 12$
- $Accuracy = 0.5$
- $Precision(Dog) \sim 0.3$
- $Precision(Cat) = 0.75$
- $Precision(Mouse) = 0.4$
- $\overline{Precision} = 0.48$

True \ Predicted	Predicted		
	Dog	Cat	Mouse
Dog	1	0	1
Cat	0	3	2
Mouse	2	1	2

- $\#Example = 12$
- $Accuracy = 0.5$
- $Precision(Dog) \sim 0.3$
- $Precision(Cat) = 0.75$
- $Precision(Mouse) = 0.4$
- $\overline{Precision} = 0.48$
- $Recall(Dog) = 0.5$
- $Recall(Cat) = 0.6$
- $Recall(Mouse) = 0.4$
- $\overline{Recall} = 0.5$

True \ Predicted			
	Dog	Cat	Mouse
Dog	1	0	1
Cat	0	3	2
Mouse	2	1	2

- $\#Example = 12$
- $Accuracy = 0.5$
- $Precision(Dog) \sim 0.3$
- $Precision(Cat) = 0.75$
- $Precision(Mouse) = 0.4$
- $\overline{Precision} = 0.48$
- $Recall(Dog) = 0.5$
- $Recall(Cat) = 0.6$
- $Recall(Mouse) = 0.4$
- $\overline{Recall} = 0.5$
- $F1Score \sim 0.49$

True \ Predicted			
	Dog	Cat	Mouse
Dog	1	0	1
Cat	0	3	2
Mouse	2	1	2

Full-Batch Learning

Full Dataset Length

$$E = \frac{1}{2m} \sum_{i=1}^m (\hat{y}_i - y_i)^2$$

repeat until convergence{

$$\theta_1 := \theta_1 - \alpha * \frac{\partial E(\theta_1)}{\partial \theta_1}$$

}

Stochastic Gradient Descent

One Example

$$E = \frac{1}{2m} \sum_{i=1}^m (\hat{y}_i - y_i)^2 = \frac{1}{2} (\hat{y}_i - y_i)^2$$

repeat until convergence{

$$\theta_1 := \theta_1 - \alpha * \frac{\partial E(\theta_1)}{\partial \theta_1}$$

}

Minibatch

Length of Minibatch

$$E = \frac{1}{2m} \sum_{i=1}^m (\hat{y}_i - y_i)^2$$


repeat until convergence{

$$\theta_1 := \theta_1 - \alpha * \frac{\partial E(\theta_1)}{\partial \theta_1}$$


}

	Stochastic GD	Full-Batch	Mini-Batch
Online-Learning	Yes	No	Yes
Variance of Gradients	High	Low	Middle
Fit in Memory	Yes	Depends on dataset length	Depends on batch-size
Speed of Convergence	High	Slow	Medium


Validation




K-Fold Cross Validation



K -Fold Cross Validation



K-Fold Cross Validation



Randomize dataset

Split whole dataset into K datasets (folds)

sum_performance = 0

for $i = 0$ in range(K){

 testset = fold(i)

 trainset=fold($i \neq 0$)

 train(trainset)

 performance = test(testset)

 sum_performance+=performance

}

sum_performance /= K

Randomize dataset

Randomize dataset

Split whole dataset into K datasets (folds)

Randomize dataset

Split whole dataset into K datasets (folds)

sum_performance = 0

for $i = 0$ in range(K){

}

Randomize dataset

Split whole dataset into K datasets (folds)

sum_performance = 0

for $i = 0$ in range(K){

 testset = fold(i)

}

Randomize dataset

Split whole dataset into K datasets (folds)

sum_performance = 0

for $i = 0$ in range(K){

 testset = fold(i)

 trainset=fold($i \neq 0$)

}

Randomize dataset

Split whole dataset into K datasets (folds)

sum_performance = 0

for $i = 0$ in range(K){

 testset = fold(i)

 trainset=fold($i \neq 0$)

 train(trainset)

 performance = test(testset)

}

Randomize dataset

Split whole dataset into K datasets (folds)

sum_performance = 0

for $i = 0$ in range(K){

 testset = fold(i)

 trainset=fold($i \neq 0$)

 train(trainset)

 performance = test(testset)

 sum_performance+=performance

}

Randomize dataset

Split whole dataset into K datasets (folds)

sum_performance = 0

for $i = 0$ in range(K){

 testset = fold(i)

 trainset=fold($i \neq 0$)

 train(trainset)

 performance = test(testset)

 sum_performance+=performance


}

sum_performance /= K

K -Fold Cross Validation

- Recommendation:
 $K=10$ (Gold Standard)
- Special Case $K = \text{\#Examples}$:
Leave One Out Cross Validation (LOOCV)

Momentum



Usual Weight Update:

$$\theta := \theta - \alpha \frac{\partial E(\theta)}{\partial \theta}, \Delta\theta = -\alpha \frac{\partial E(\theta)}{\partial \theta}$$

$$\theta := \theta + \Delta\theta$$

Usual Weight Update:

$$\theta := \theta - \alpha \frac{\partial E(\theta)}{\partial \theta}, \Delta\theta = -\alpha \frac{\partial E(\theta)}{\partial \theta}$$

$$\theta := \theta + \Delta\theta$$

Momentum:

$$\Delta\theta_t := -\alpha \frac{\partial E(\theta)}{\partial \theta} + \varepsilon * \Delta\theta_{t-1}, 0 < \varepsilon \leq 1$$


$$\theta_{t+1} = \theta_t + \Delta\theta_t$$

Momentum:

- **Faster convergence**
- **Smooth out variance of gradients**
- **Use Decay for Momentum ε**
- **Find good value and decay through cross validation**

Rate performance of neural net

- High bias / Underfitting
- High variance / Overfitting



Dropout

- **Strategy to avoid overfitting**
- **Randomly exclude some neurons from training process**
 - **Exclude from forward pass**
 - **Exclude from backward pass**

tensorboard

- **Create a tf.summary**
 - Store important variables (Scalars, Histograms, Images, ...)
- **Write tf.summary to file**
- **Open tensorboard with events-file as target**

tensorboard examples:

- **Tensorflow documentation**
- **Internet**
- **Star Recognition (Repository)**
- **DQN (Repository)**



CNN example

- **Start tensorboard to debug the training process**
- **Implement Momentum**
- **Use a data augmentation technique to make generalization more robust**
- **Implement train, test, validation split**
- **Download weights and image of one of the five stars to test prediction**

Start tensorboard to debug the training process

- Start the training process
- Start tensorboard in directory with event-file

Hints: Implement Momentum

- Implement `tf.train.MomentumOptimizer`
- Search for „optimizer“ in `cnn.py`

Use a data augmentation technique to make generalization more robust

- Change .tfrecord writing to save RGB-Images
- Implement RGB to Grayscale in training process
- Normalize data before the training process!!!
- Use `tf.image.random...` operations to augment the input data

Implement train, test, validation split

- Create a new python script
- Count the number of examples
- Randomize the examples
- Create variables for the train, test and validation proportion
- Write .tfrecord files
 - Have a look at `write_tfrecord.py`

Download weights and image of one of the five stars to test prediction

- Works only if RGB to Grayscale conversion is implemented
- Solution will be pushed at the end of the lesson

- **Data (if not yet downloaded):**
<https://nextcloud.mirevi.medien.hs-duesseldorf.de/index.php/s/kPXwJiac7vTQVeu>
- **Pretrained Weights:**
<https://nextcloud.mirevi.medien.hs-duesseldorf.de/index.php/s/L6Y6tnD3PpANKmr>
- **Repository:** <https://github.com/mati3230/modalg181>
- **Read:** https://www.tensorflow.org/tutorials/deep_cnn